

01_cleaning_walkthrough

June 24, 2025

```
[1]: # Step 1: Import the libraries we'll use
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, StandardScaler

# This makes sure we can see all the columns
pd.set_option('display.max_columns', None)

# Step 2: Load the Titanic CSV from the data folder
df = pd.read_csv('../data/Titanic-Dataset.csv') # if your file was named train.
↳ csv, use that instead
df.head()
```

Matplotlib is building the font cache; this may take a moment.

```
[1]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[2]: # Step 3: Let's check for missing values and data types
print(" Data Types:\n", df.dtypes)
print("\n Missing values in each column:")
print(df.isnull().sum())
```

```
Data Types:
PassengerId      int64
Survived          int64
Pclass           int64
Name             object
Sex              object
Age             float64
SibSp            int64
Parch            int64
Ticket           object
Fare             float64
Cabin            object
Embarked         object
dtype: object
```

```
Missing values in each column:
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64
```

```
[3]: # Step 4: Fixing missing values

# 1. Fill missing Age with the median
df['Age'].fillna(df['Age'].median(), inplace=True)

# 2. Fill missing Embarked with the mode (most frequent)
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)

# 3. Drop the Cabin column entirely
df.drop(columns='Cabin', inplace=True)

# Check again to confirm all missing values are gone
df.isnull().sum()
```

C:\Users\dutta\AppData\Local\Temp\ipykernel_8556\3291964984.py:4: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df['Age'].fillna(df['Age'].median(), inplace=True)
```

C:\Users\dutta\AppData\Local\Temp\ipykernel_8556\3291964984.py:7: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)
instead, to perform the operation inplace on the original object.

```
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
```

```
[3]: PassengerId    0  
     Survived      0  
     Pclass       0  
     Name         0  
     Sex          0  
     Age          0  
     SibSp        0  
     Parch        0  
     Ticket       0  
     Fare         0  
     Embarked     0  
     dtype: int64
```

```
[4]: # Step 5: Convert text columns into numbers  
  
from sklearn.preprocessing import LabelEncoder  
  
# Create encoder  
le = LabelEncoder()
```

```

# Encode 'Sex' and 'Embarked'
df['Sex'] = le.fit_transform(df['Sex'])
df['Embarked'] = le.fit_transform(df['Embarked'])

# Drop 'Name' and 'Ticket' columns
df.drop(columns=['Name', 'Ticket'], inplace=True)

# Display first few rows to see changes
df.head()

```

```

[4]:
  PassengerId  Survived  Pclass  Sex  Age  SibSp  Parch    Fare  Embarked
0           1         0       3    1  22.0     1     0    7.2500         2
1           2         1       1    0  38.0     1     0   71.2833         0
2           3         1       3    0  26.0     0     0    7.9250         2
3           4         1       1    0  35.0     1     0   53.1000         2
4           5         0       3    1  35.0     0     0    8.0500         2

```

```

[5]: # Step 6: Standardize Age and Fare so they're on the same scale
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# Fit the scaler on Age and Fare, then transform them
df[['Age', 'Fare']] = scaler.fit_transform(df[['Age', 'Fare']])

# Quick sanity-check: the new mean should be ~0 and std ~1
df[['Age', 'Fare']].describe().loc[['mean', 'std']]

```

```

[5]:
      Age      Fare
mean  2.272780e-16  3.987333e-18
std   1.000562e+00  1.000562e+00

```

```

[6]: # Step 7: Visualise and trim Fare outliers

import seaborn as sns
import matplotlib.pyplot as plt
from pathlib import Path

# 1 Boxplot BEFORE trimming
sns.boxplot(x=df['Fare']).set_title('Fare - before trimming')
plt.tight_layout()

# Save the picture to images/ so you can show it in GitHub
Path('../images').mkdir(exist_ok=True)
plt.savefig('../images/fare_boxplot_before.png')
plt.show()

```

```

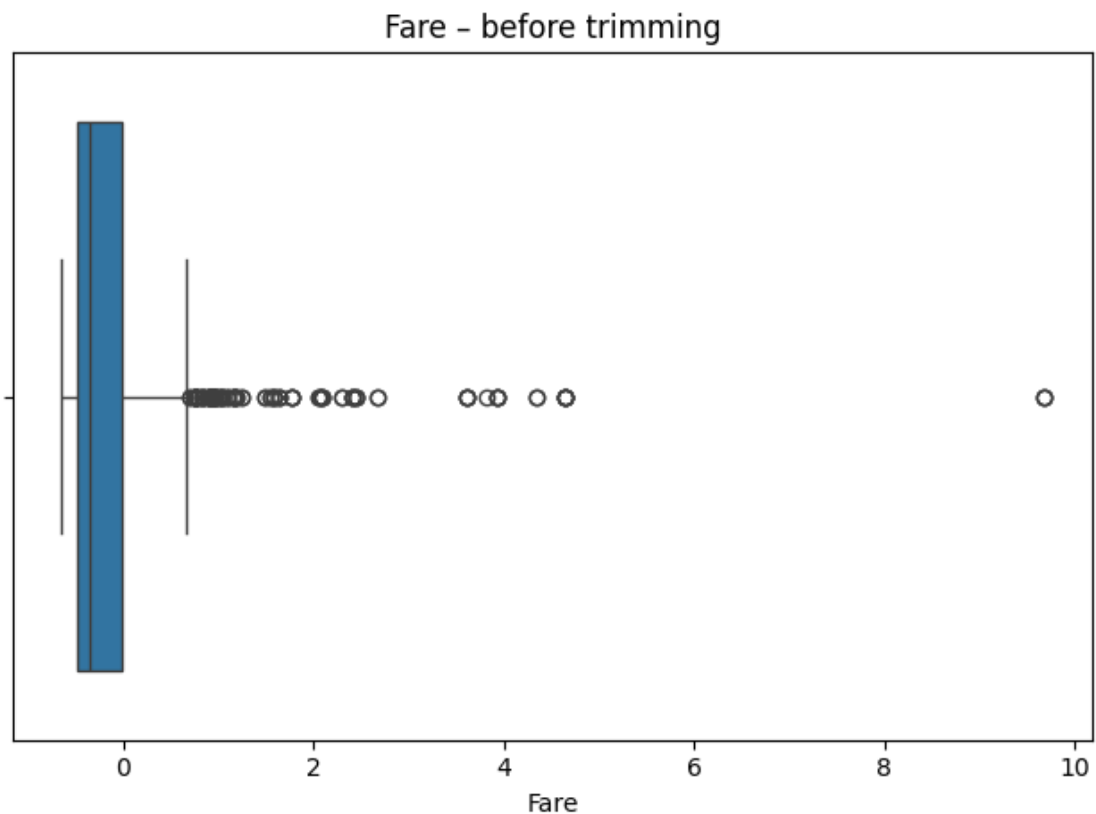
# 2 IQR method to define "too high" or "too low"
q1, q3 = df['Fare'].quantile([0.25, 0.75])
iqr = q3 - q1
lower, upper = q1 - 1.5 * iqr, q3 + 1.5 * iqr

# 3 Keep only rows within the bounds
before = len(df)
df = df[(df['Fare'] >= lower) & (df['Fare'] <= upper)]
after = len(df)

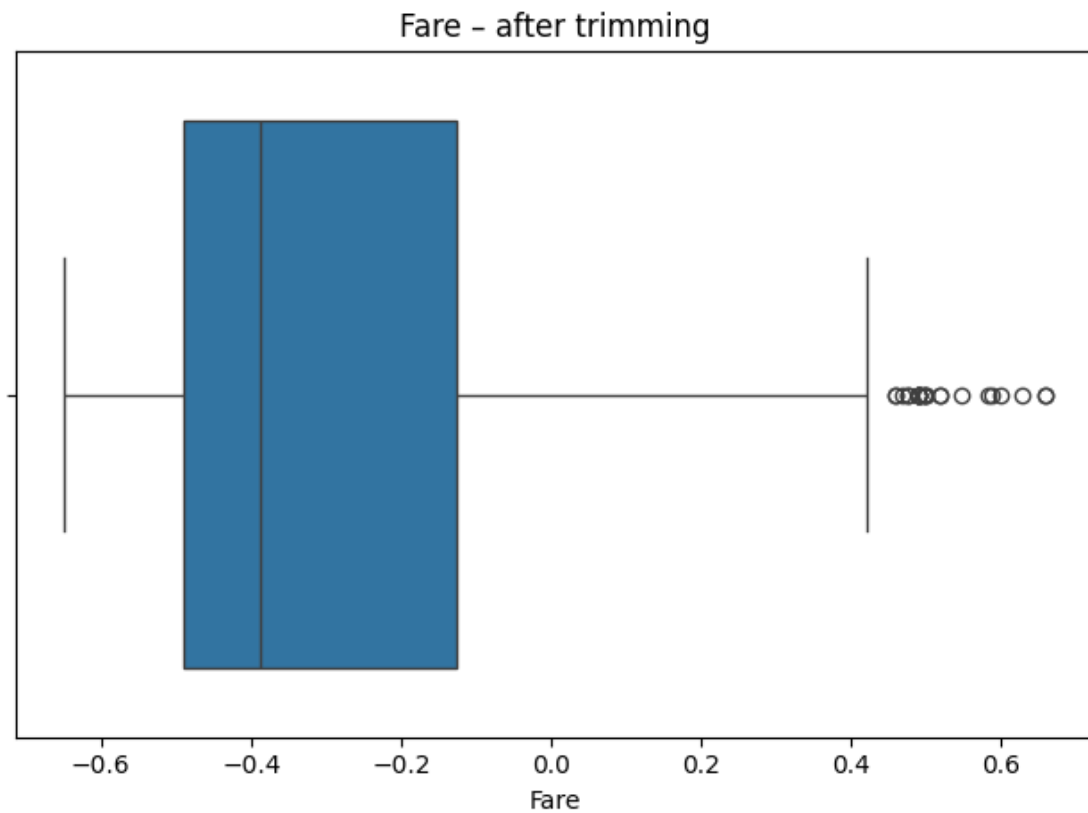
print(f"Removed {before - after} outlier rows. Rows left: {after}")

# 4 Boxplot AFTER trimming (quick check)
sns.boxplot(x=df['Fare']).set_title('Fare - after trimming')
plt.tight_layout()
plt.savefig('../images/fare_boxplot_after.png')
plt.show()

```



Removed 116 outlier rows. Rows left: 775



```
[7]: # Step 8: Save the final cleaned dataset
df.to_csv('../data/titanic_cleaned.csv', index=False)
print(f"Cleand data written to data/titanic_cleaned.csv    (rows: {len(df)})")
```

Cleand data written to data/titanic_cleaned.csv (rows: 775)

```
[ ]:
```